

Foundational Knowledge for AI

Fredrik Heintz

Dept. of Computer Science
Linköping University

fredrik.heintz@liu.se

@FredrikHeintz

Outline:

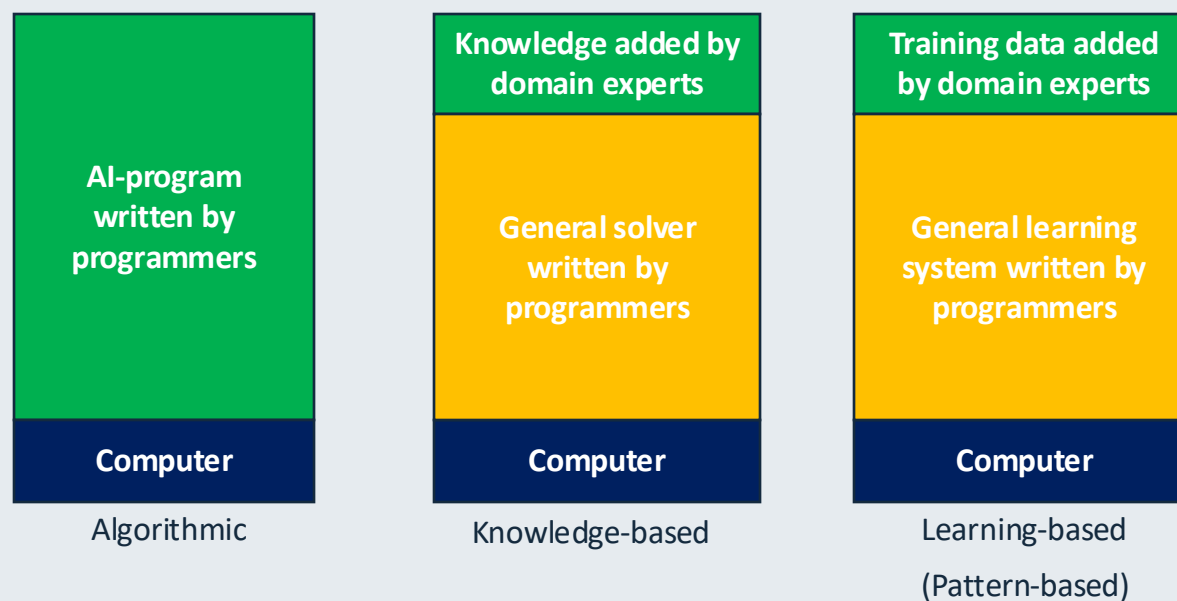
- Uninformed Search Algorithms
- Informed Search Algorithms
- Search in Complex Environments
- Adversarial Search and Games



Co-funded by
the European Union

Algorithmic, Knowledge-Based and Learning-Based AI

2

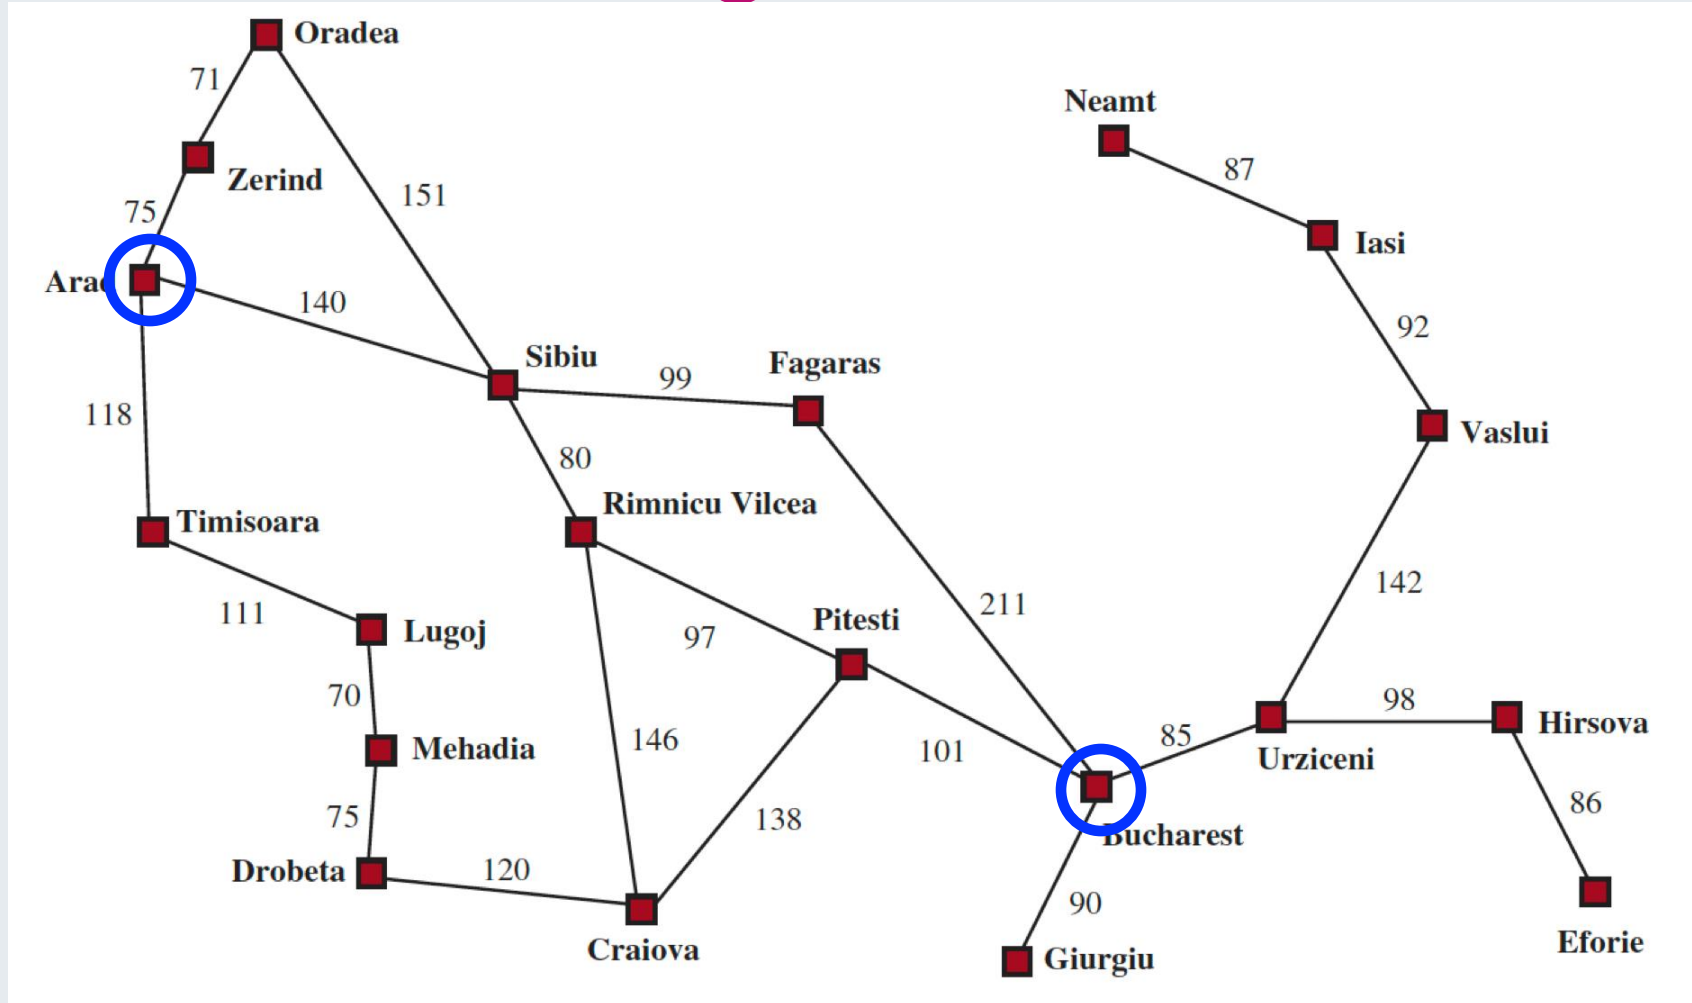


TDDC17 - HT23 - Fredrik Heintz -
LE2 Search I (based on slides by
Patrick Doherty)




Co-funded by
the European Union

Romania Route Finding Problem



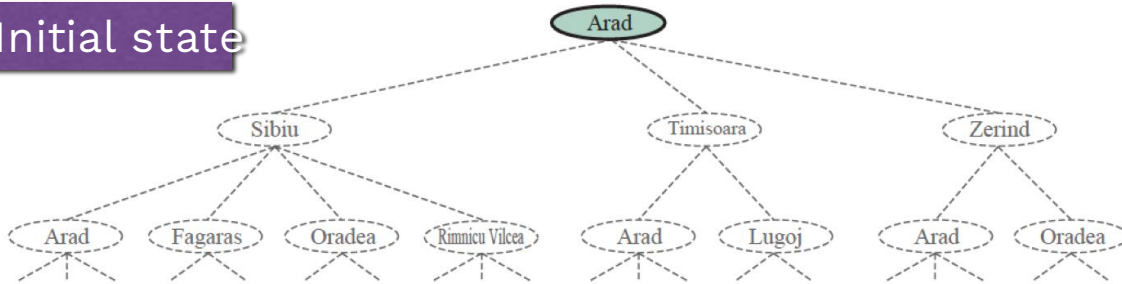
Problem Formulation

- State Space - A set of possible states that the environment can be in
 - Cities in the Romania map.
- Initial State - The state the agent starts in
 - Arad
- ACTIONS(State) - A description of what actions are applicable in each state.
 - **ACTIONS(Arad)** = {ToSibiu, ToTimisoara, ToZerind}
- RESULT(State, Action) - A description of what each action does (Transition Model)
 - **RESULT(Arad, ToZerind)** = Zerind
- Goal Test - Tests whether a given state is a goal
 - Often a set of states: { Bucharest }
- An Action Cost Function - denoted **ACTION-COST**(s, a, s') when programming, or $c(s, a, s')$ when doing math.
 - Gives the numeric cost of doing a in s to reach state s' .
 - Cost functions should reflect the agents performance measure: distance, time taken, etc.
- Solution - A path from the start state to the goal state

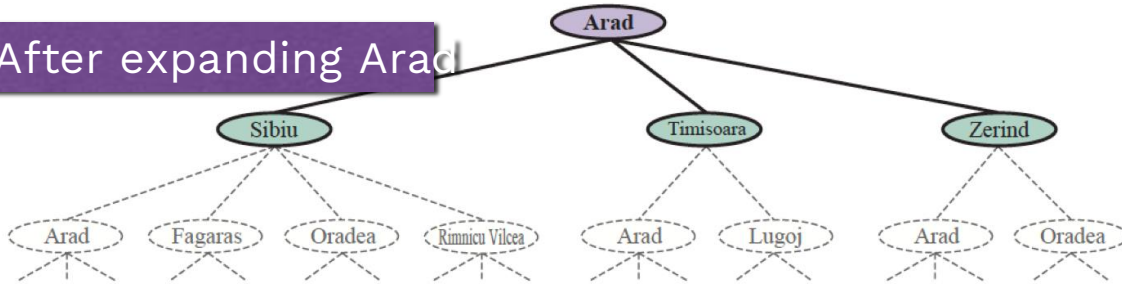


3 Partial Search Trees: Route Finding

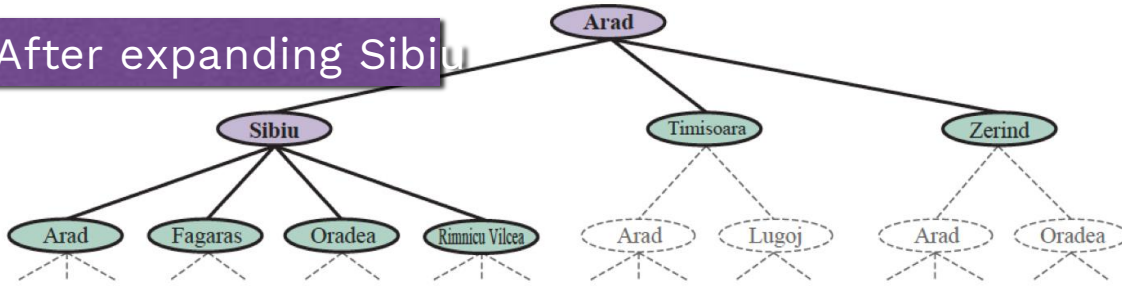
Initial state



After expanding Arad



After expanding Sibiu



Expanded

Generated

Not Expanded






Figure 3.1 A simplified road map of part of Romania, with road distances in miles.

Problem Formulation: Vacuum cleaning world

- **State Space:**
2 positions, dirt or no dirt
→ 8 world states
- **Actions:**
Left (L), Right (R), or Suck (S)
- **Transition model:** next slide
- **Initial State:** Choose.
- **Goal States:**
States with no dirt in the rooms
- **Action costs:**
one unit per action



Solving the Vacuum World



Problem Formulation: Missionaries and Cannibals

Informal problem description:

- Three missionaries and three cannibals are on **one side** of a river that they wish to cross.
- A boat is available that can hold **at most two people**.
- You must never leave a group of **missionaries outnumbered by cannibals** on the same bank.
 - How should the **state space** be represented?
 - What is the **initial state**?
 - What is the **goal state**?
 - What are the **actions**?

One Formalization

- Many other formalisations

State Space: triple (x,y,z) with $0 \leq x,y,z \leq 3$, where x,y , and z represent the number of missionaries, cannibals and boats currently on the original bank.

Initial State: $(3,3,1)$

Actions: see transition model

Transition Model: from each state, either bring one missionary, one cannibal, two missionaries, two cannibals, or one of each type to the other bank.

Note: not all states are attainable (e.g., $(0,0,1)$), and some are illegal.

Goal States: $\{(0,0,0)\}$



Action Costs: 1 unit per crossing

General Search

From the initial state, produce all successive states step by step → search tree.

M, C, B

(a) initial state (3,3,1)



Examples of Real-World Problems

- Route Planning, Shortest Path Problem
 - Routing video streams in computer networks, airline travel planning, military operations planning...
- Travelling Salesperson Problem (TSP)
 - A common prototype for NP-complete problems
- VLSI (integrated circuits) Layout
 - Another NP-complete problem
- Robot Navigation (with high degrees of freedom)
 - Difficulty increases quickly with the number of

degrees of freedom. Further possible complications: errors of perception, unknown environments

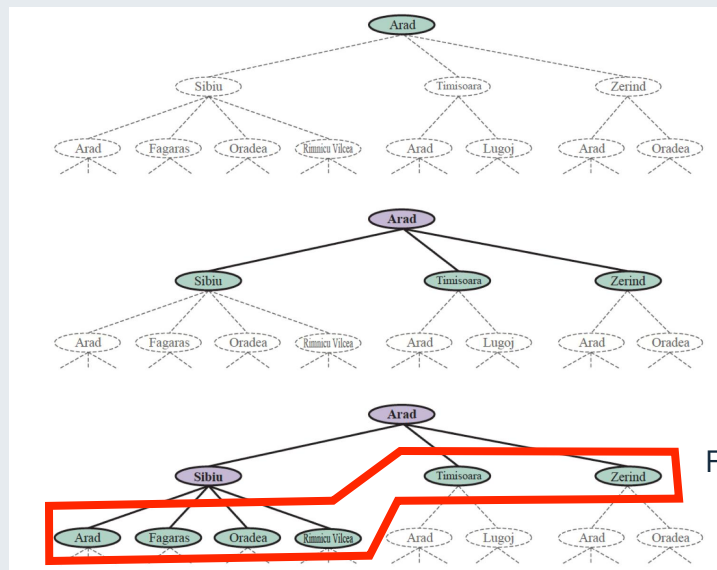
- Assembly Sequencing
 - Planning of the assembly of complex objects (by robots)

Search Algorithms

We focus on algorithms that superimpose a search tree over the state space graph

Important distinction between search tree and state space!

Search tree Construction



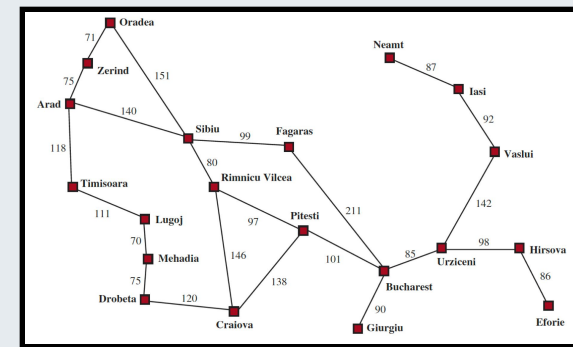
Expanded

Generated

Not Expanded

Frontier (choose from here)


State Space Graph



Graph-Search: Only add a child if the state associated with it has not already been reached:

- Avoid cycles
- Avoid redundant paths


Oradea: has 2 successor states, but already reached by other paths, so do not expand.



Sequence of search trees superimposed
(each node only has one parent)

Search Algorithms

Separation property of graph search



The frontier (green) separates the interior (lavender) from the exterior (faint dashed)

Any state with a node generated for it is said to be reached

What is the best way to search through a state space in a systematic manner in order to reach a goal state?

Search Strategies

A *strategy* is defined by picking the order of node expansion.

Strategies can be *evaluated* along the following dimensions:

Completeness – does it find a solution if it exists?

Time Complexity – number of nodes generated/expanded

Space Complexity – maximum number of nodes in memory

Optimality – does it always find a least cost solution

Time and space complexity are *measured* in terms of:

b – *maximum branching factor* of search tree

d – *depth of the least cost solution* in the search tree

m – *maximum length* of any path in the state space (possibly infinite)

Some Search Classes

- *Uninformed Search* (Blind Search)
 - No additional information about states besides that in the problem definition
 - Can only generate successors and compare against state.
 - Some examples:
 - Breadth-first search, Depth-first search, Iterative deepening DFS
- *Informed Search* (Heuristic Search)
 - Strategies have additional information as to whether non-goal states are more promising than others.
 - Some examples:
 - Greedy Best-First Search, A* Search



Implementing the Search Tree

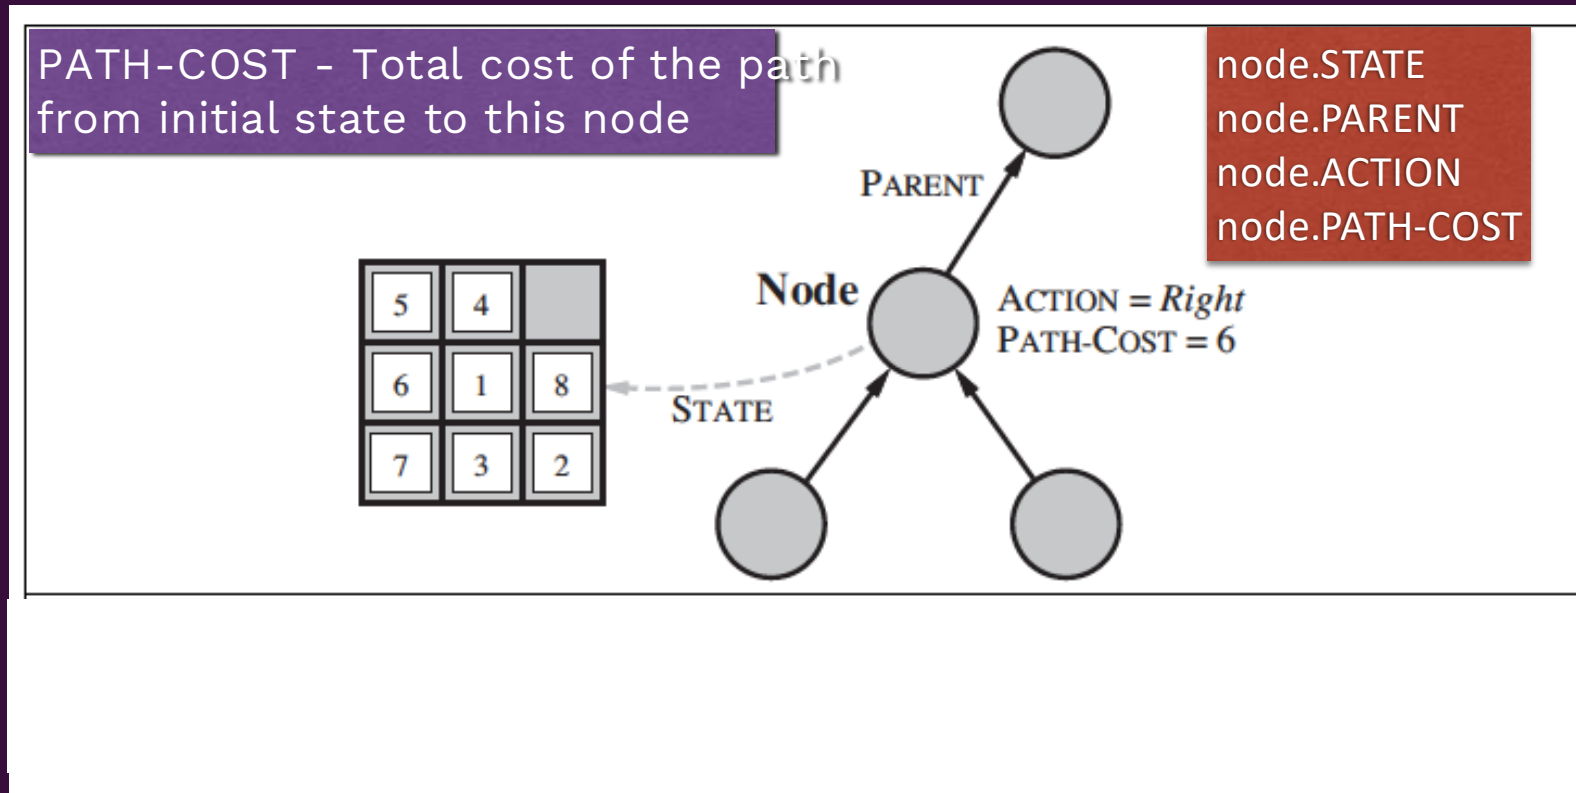
Three kinds of Queues:

- Priority Queue - min cost first
- FIFO Queue - first in, first out
- LIFO Queue - last in, first out

- Data structure for nodes in the search tree:
- STATE: state in the state space
- PARENT: Predecessor nodes
- ACTION: The operator that generated the node
- DEPTH: number of steps along the path from the initial state
- PATH-COST: Cost of the path from the initial state to the node
- Operations on a queue/frontier (4th Edition):
- IS-EMPTY(*frontier*): Empty test
- POP(*frontier*): Returns the first element of the queue
- TOP(*frontier*): Returns the first element
- ADD(*node*, *frontier*): Inserts new elements into the queue
- Operations on a queue (3rd Edition):
- Make-Queue(Elements): Creates a queue
- Empty?(Queue): Empty test
- First(Queue): Returns the first element of the queue
- Remove-First(Queue): Returns the first element
- Insert(Element, Queue): Inserts new elements into the queue
- (various possibilities)
- Insert-All(Elements, Queue): Inserts a set of elements into the queue



States (in state space) and Nodes (in a search tree)



Can have several nodes with the same state due to multiple paths to the state

The search tree describes paths between states leading towards a goal

Breadth-First Search

Assume all actions have the same cost

Search by Minimal Depth:




Figure 3.8 Breadth-first search on a simple binary tree. At each stage, the node to be expanded next is indicated by the triangular marker.

Could be implemented using Best-First-Search:

```
function | BREADTH-FIRST-SEARCH(problem) returns a solution node, or failure  
return BEST-FIRST-SEARCH(problem, DEPTH )
```

$$f(n) = n.DEPTH$$

Is a better way!

Depth-First Search

Assume all actions have the same cost

Always expands
deepest node
in the frontier first

Usually implemented not as
graph search but as
tree-like search (without a
table of reached states)





Figure 3.11 A dozen steps (left to right, top to bottom) in the progress of a depth-first search on a binary tree from start state A to goal M. The frontier is in green, with a triangle marking the node to be expanded next. Previously expanded nodes are lavender, and potential future nodes have faint dashed lines. Expanded nodes with no descendants in the frontier (very faint lines) can be discarded.

Heuristic Search

Intuitions behind Heuristic Search



Find a heuristic measure $h(n)$ which estimates how close a node n in the frontier is to the nearest goal state and then order the frontier queue accordingly relative to closeness.

The evaluation function $f(n)$, previously discussed will include $h(n)$:

$$f(n) = \dots + h(n)$$

$h(n)$ is intended to provide domain specific hints about location of goals

Romania Travel Problem

Let's find a heuristic!

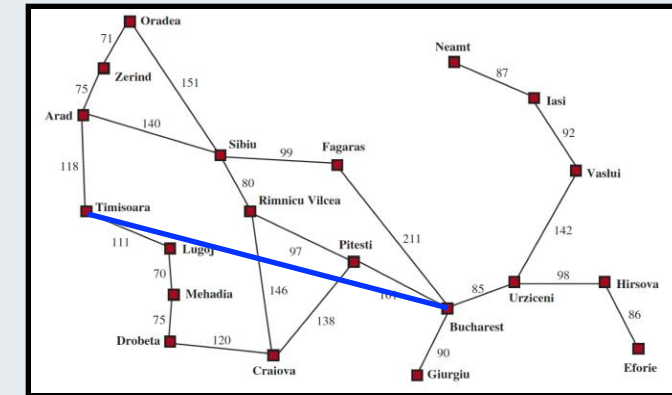
Straight line distance from city n to goal city n'

Assume the cost to get somewhere is a function of the distance traveled

Straight line distance to Bucharest from any city

$h_{SLD}()$

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

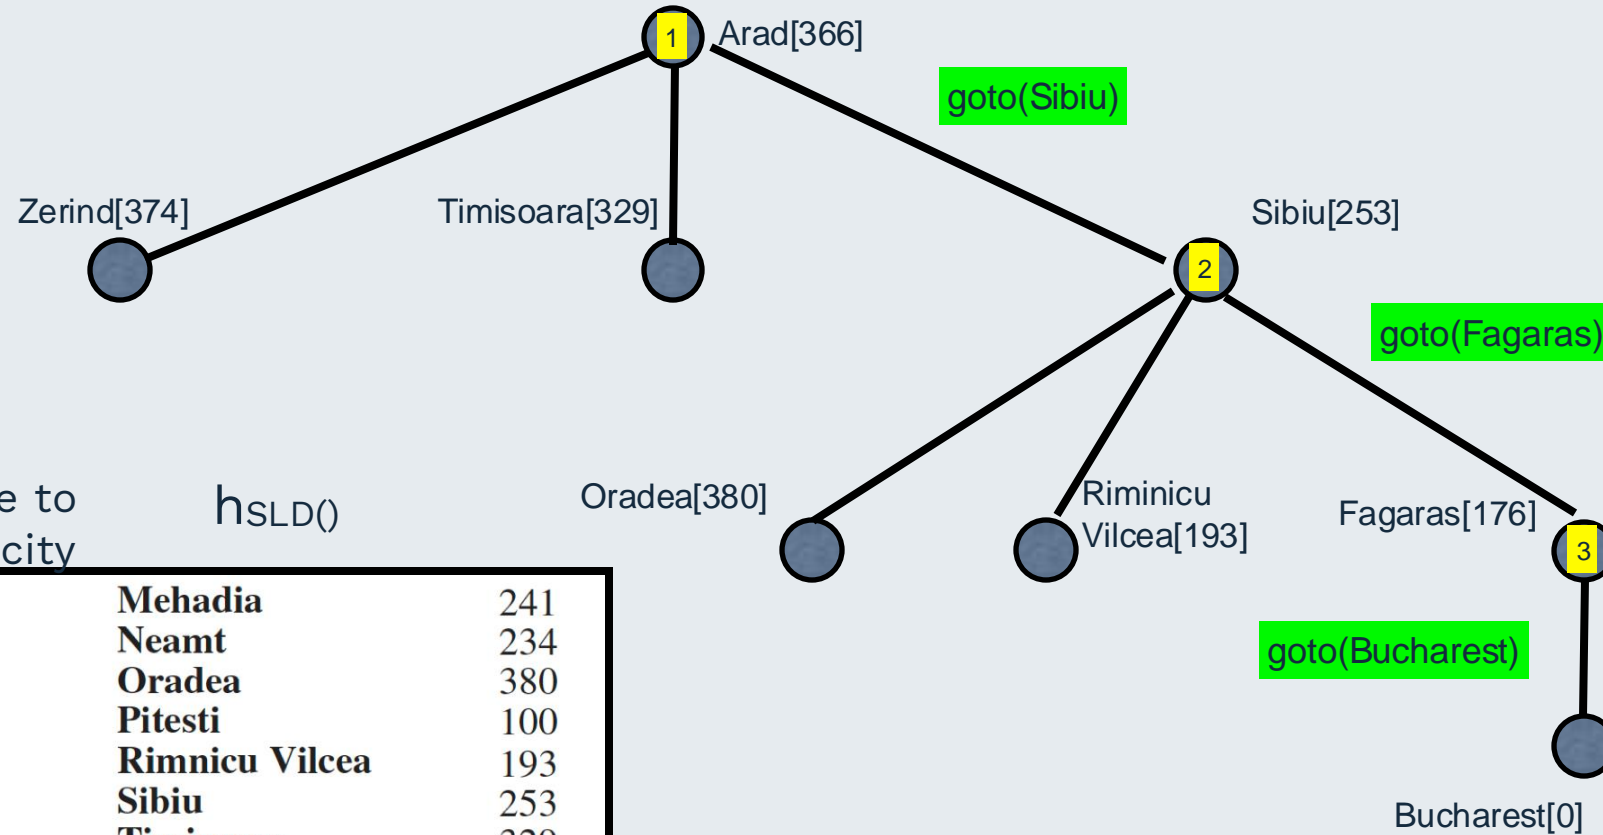


Heuristic:

$$f(n) = h_{SLD}(n)$$

Notice the SLD under estimates the actual cost!

Greedy Best-First Search: Romania



Straight line distance to
Bucharest from any city

$h_{SLD}()$

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

A*-1

(a) The initial state



$$366 = 0 + 366$$

$$g(\text{Arad}) + h(\text{Arad})$$

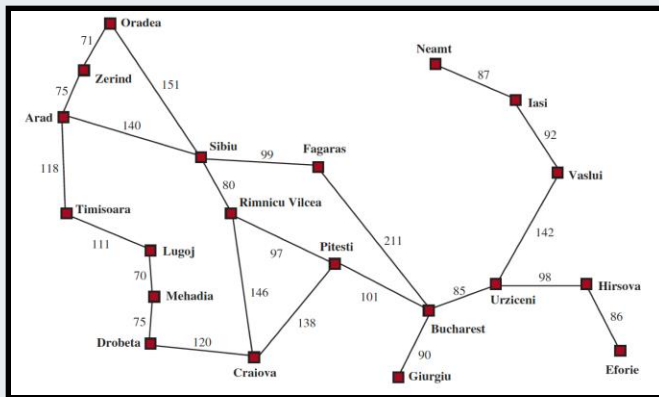
Heuristic (with Bucharest as goal):

$$f(n) = g(n) + h(n)$$

$g(n)$ - Actual distance from root node to n

$h(n)$ - $h_{SLD}(n)$ straight line distance from n to Bucharest

$g(n)$




Straight line distance to Bucharest from any city

$$h(n) = h_{SLD}(n)$$


Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

A*-2

(b) After expanding Arad



$g(n)$

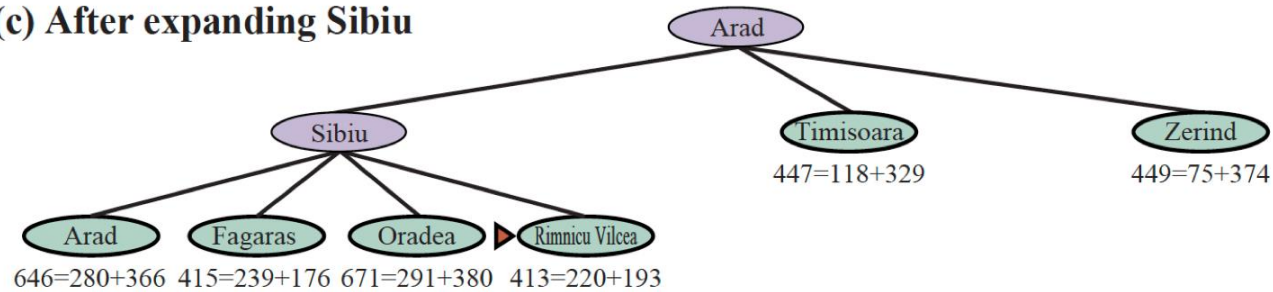


Straight line distance to Bucharest from any city $h(n) = h_{SLD}(n)$

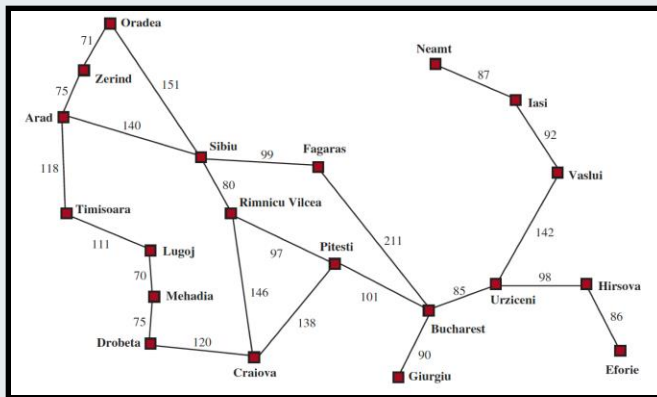
Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

A*-3

(c) After expanding Sibiu



$g(n)$

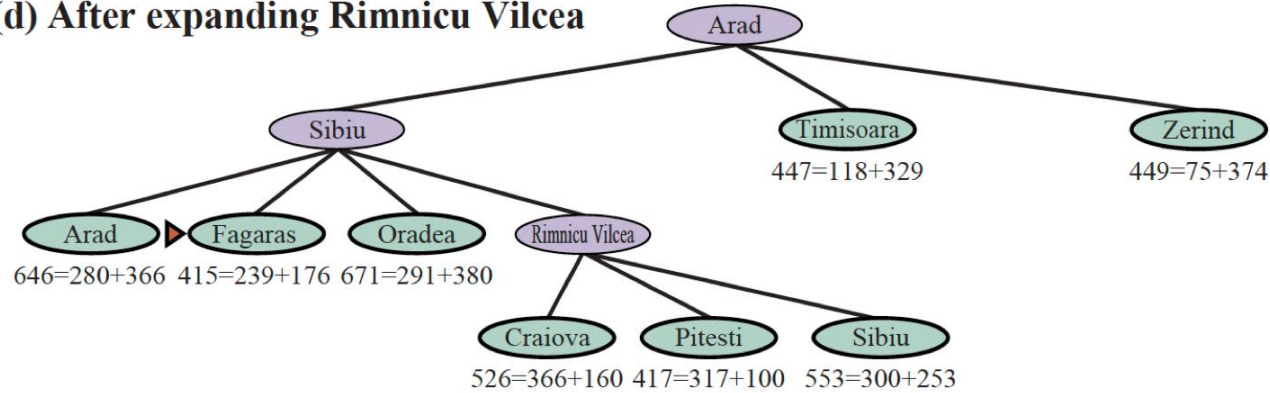


Straight line distance to Bucharest from any city $h(n) = h_{SLD}(n)$


Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

A*-4

(d) After expanding Rimnicu Vilcea



$g(n)$

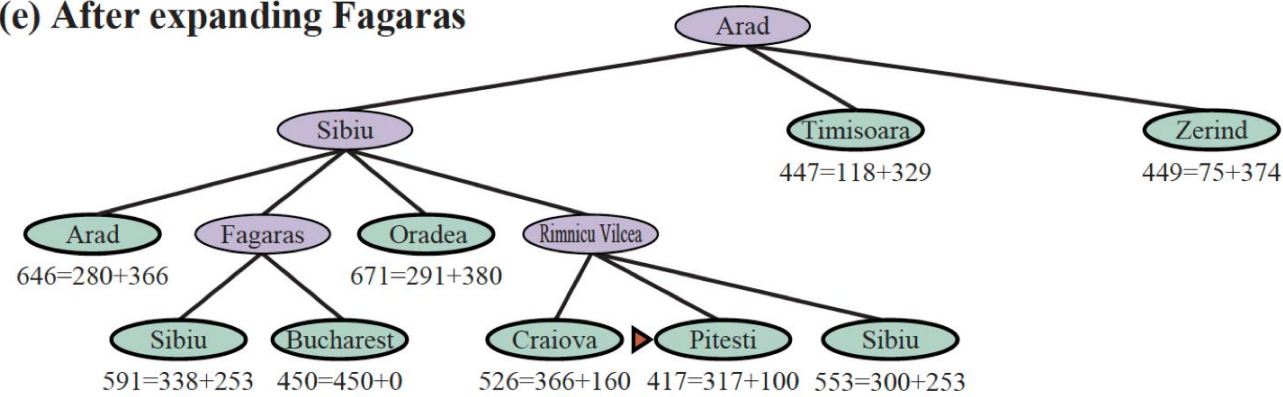


Straight line distance to Bucharest from any city $h(n) = h_{SLD}(n)$

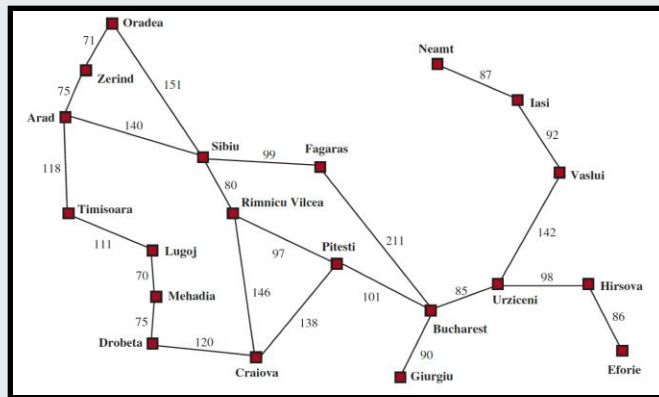
Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

A*-4

(e) After expanding Fagaras



$g(n)$

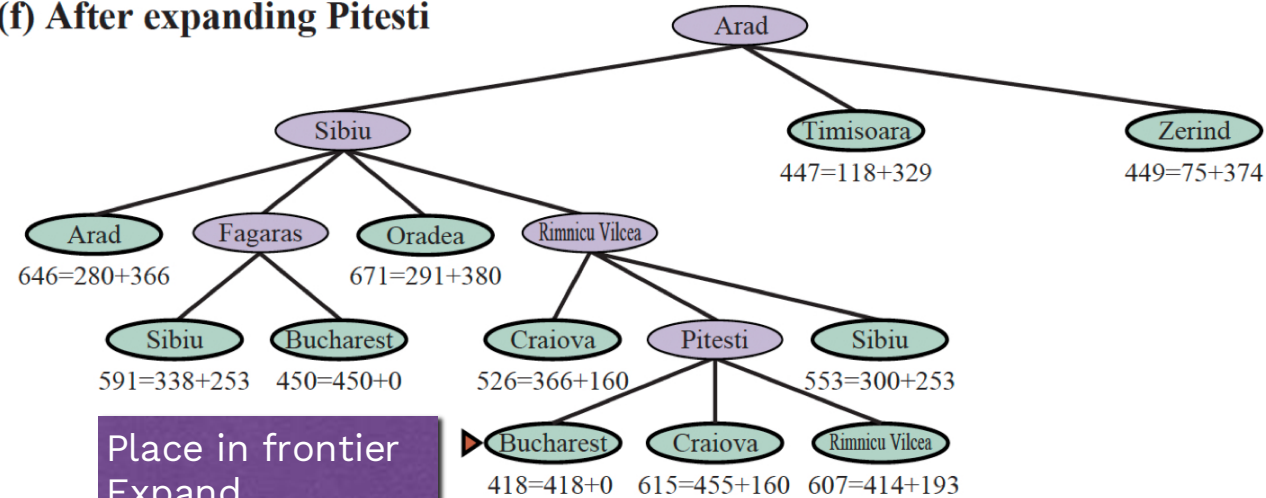


Straight line distance to Bucharest from any city $h(n) = h_{SLD}(n)$

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

A*-6

(f) After expanding Pitesti

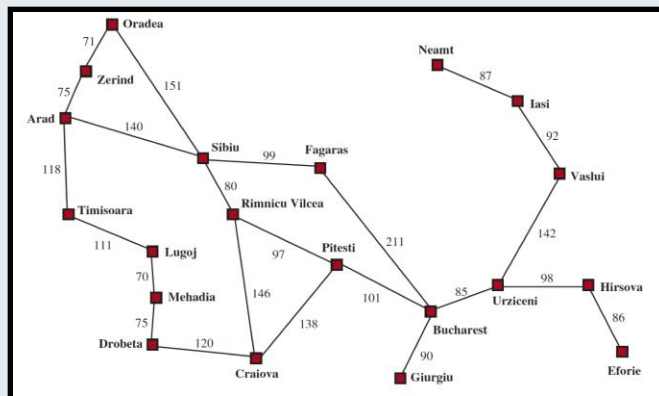


Place in frontier
Expand
Late Testing Goal

$g(n)$

Straight line distance to
Bucharest from any city

$h(n) = h_{SLD}(n)$



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Some properties of A*


- Cost-Optimal -
 - for a given admissible heuristic (tree-like search)
 - for a given consistent heuristic (tree-like, graph-search)
 - Consistent heuristics are admissible heuristics but not vice-versa.
- Complete - Eventually reach a contour equal to the path of the least-cost to the goal state.
- Optimally efficient - No other algorithm, that extends search paths from a root is guaranteed to expand fewer nodes than A* for a given heuristic function.
- The exponential growth for most practical heuristics will eventually overtake the computer (run out of memory)
 - The number of states within the goal contour is still exponential in the length of the solution.
 - There are variations of A* that bound memory....




Search in Complex Environments



Local Search: 8 Queens Problem



Bad Solution



Good Solution

Problem:

Place 8 queens on a chessboard such that
No queens attacks another

- Local Search:

- the path to the goal is irrelevant!
- we do not care about reached states
- complete state formulation is a straightforward representation:
 - 8 queens, one in each column
- operate by searching from start state to neighbouring states, choose the best neighbour so far, repeat

8 Queens is a candidate for use of local search!

8^8 (about 16 million configurations)

Local Search Techniques

- Advantages:
 - They use very little memory
 - Often find solutions in large/infinite search spaces where systematic algorithms would be unreasonable
 - Can be used to solve optimisation problems
- Disadvantages
 - Since they are not systematic they may not find solutions because they leave parts of the search space unexplored.
 - Performance is dependent on the topology of the search space
 - Search may get stuck in local optima

Global Optimum: The best possible solution to a problem.

Local Optimum: A solution to a problem that is better than all other solutions that are slightly different, but worse than the global optimum

Greedy Local Search: A search algorithm that always takes the best immediate, or local, solution while finding an answer. Greedy algorithms find the overall, or globally optimal solution for some optimization problems, but may find less-than-optimal solutions for some instances of other problems. (They may also get stuck!)


Hill-Climbing Algorithm (steepest ascent version)

```
function HILL-CLIMBING(problem) returns a state that is a local maximum  
  current  $\leftarrow$  problem.INITIAL  
  while true do  
    neighbor  $\leftarrow$  a highest-valued successor state of current  
    if VALUE(neighbor)  $\leq$  VALUE(current) then return current  
    current  $\leftarrow$  neighbor
```

When using heuristic functions: steepest descent version

Greedy Progress: Hill Climbing

Aim: Find the global maximum




Hill Climbing: Modify the current state to try and improve it

Variations on Hill Climbing


- [Stochastic Hill Climbing](#)
 - Choose among uphill moves at random, weighting choice by probability with the steepness of the move
- [First Choice Hill Climbing](#)
 - Implements stochastic hill climbing by randomly generating successors until one is generated that is better than the current state.
- [Random-Restart Hill Climbing](#)
 - Conducts a series of hill-climbing searches from randomly generated initial states until a goal is found.

Local Beam Search

Start with k
random states

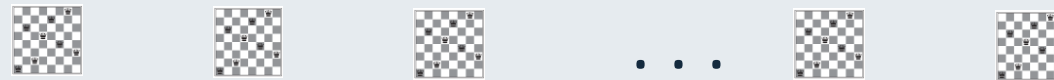


Determine
successors
of all k
random states



If any successors are goal states
then finished

Else select k
best states from
union of
successors and
repeat



Can suffer from lack of diversity among the k states
(concentrated in small region of search space).

Stochastic variant: choose k successors at random with
probability of choosing the successor being an increasing
function of its value.

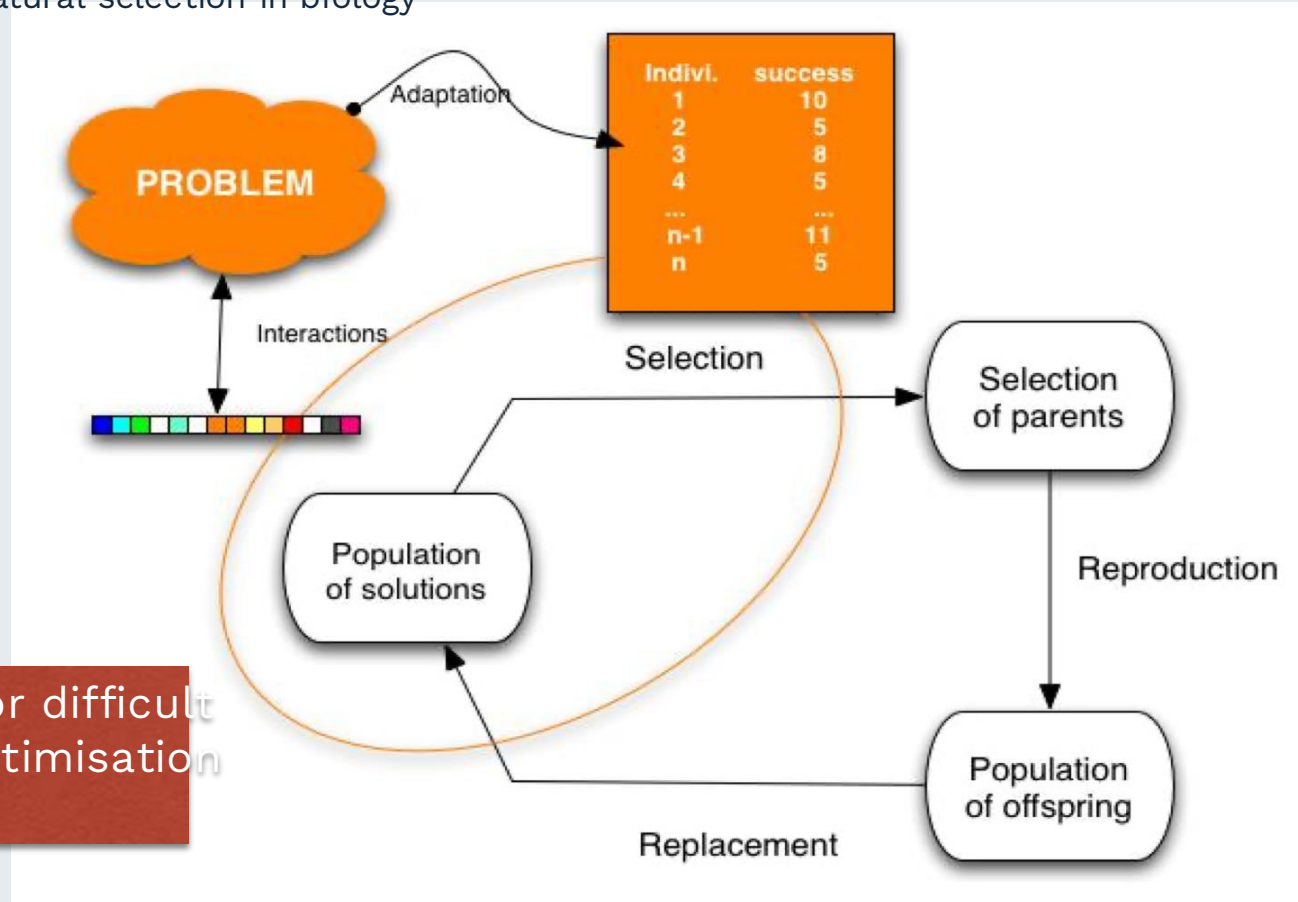
Simulated Annealing

Hill Climbing + Random Walk

- Escape local maxima by allowing “bad” moves (random)
 - **Idea:** but gradually decrease their size and frequency
 - Origin of concept: metallurgical annealing
- Bouncing ball analogy (gradient descent):
 - Shaking hard (= high temperature)
 - Shaking less (= lower the temperature)
- If Temp decreases slowly enough, best state is reached


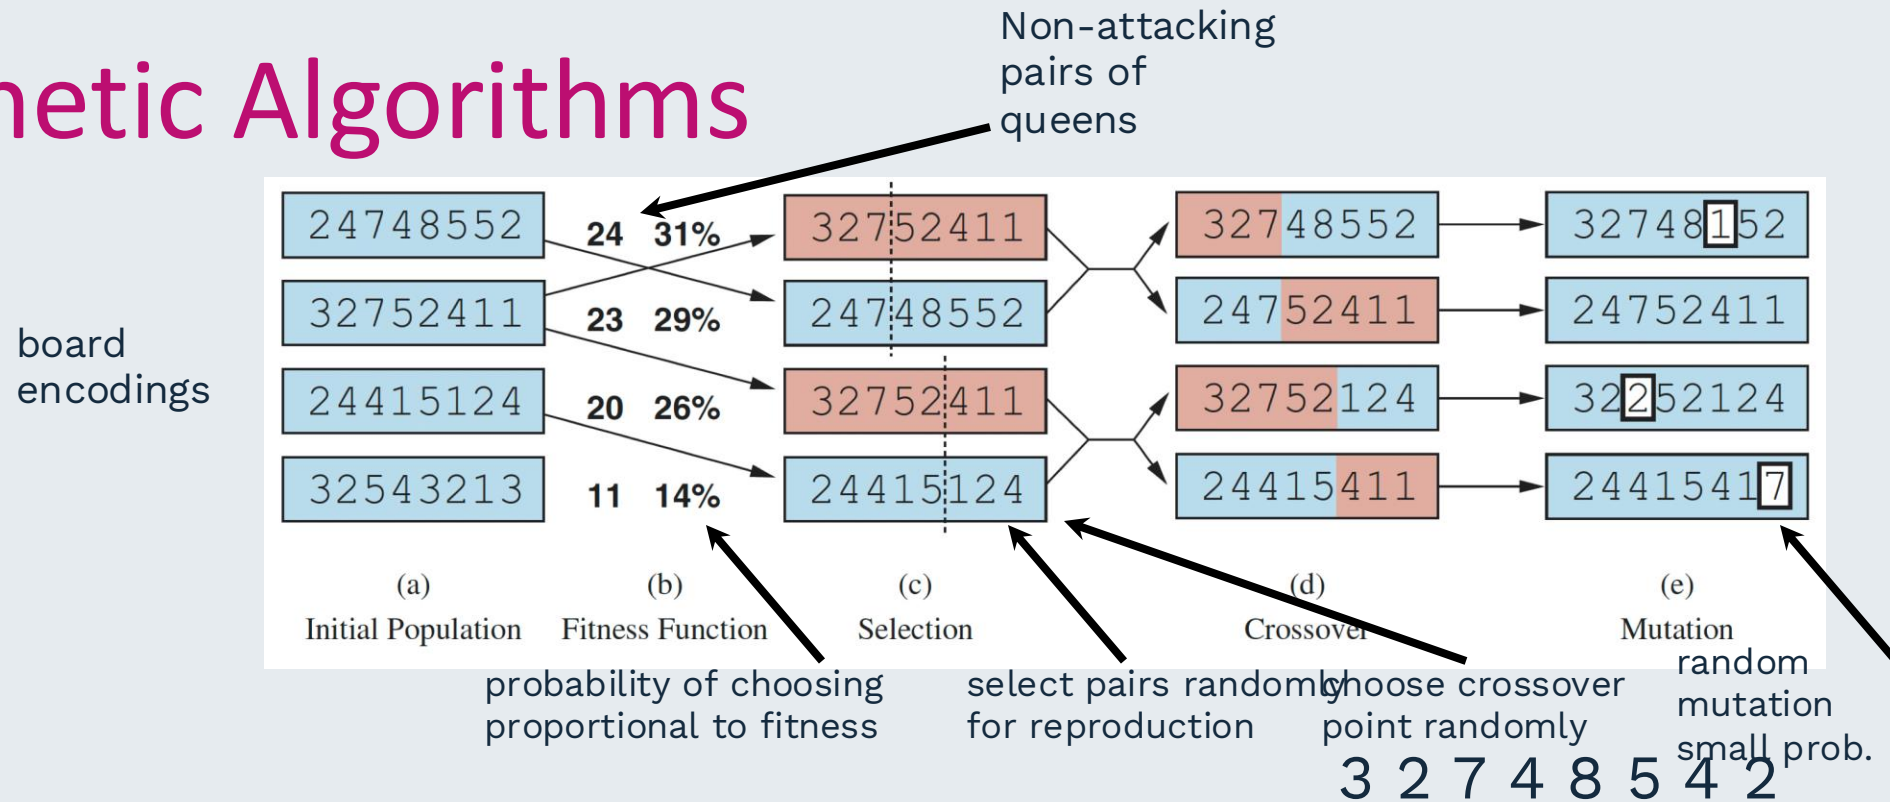
Evolutionary Algorithms

Variants of Stochastic Beam Search using the metaphor of natural selection in biology



Often used for difficult
non-linear optimisation
problems

Genetic Algorithms



Adversarial Search



Why Study Board Games?

Board games are one of the **oldest branches** of AI (Shannon and Turing 1950).

- Board games present a very abstract and **pure form** of competition between two opponents and clearly require a form of “intelligence”.
- The states of a game are **easy to represent**
- The possible **actions** of the players are well-defined
 - Realization of the game as a **search problem**
 - It is nonetheless a **contingency problem**, because the characteristics of the opponent are not known in advance

Challenges

Board games are not only difficult because they are contingency problems, but also because the search trees can become astronomically large.

Examples:

- Chess: On average 35 possible actions from every position, 100 possible moves/ply (50 each player): $35^{100} \approx 10^{150}$ nodes in the search tree (with “only” 10^{40} distinct chess positions (states)).
- Go: On average 200 possible actions with circa 300 moves: $200^{300} \approx 10^{700}$ nodes.

Good game programs have the properties that they

- delete irrelevant branches of the game tree,
- use good evaluation functions for in-between states, and
- look ahead as many moves as possible.



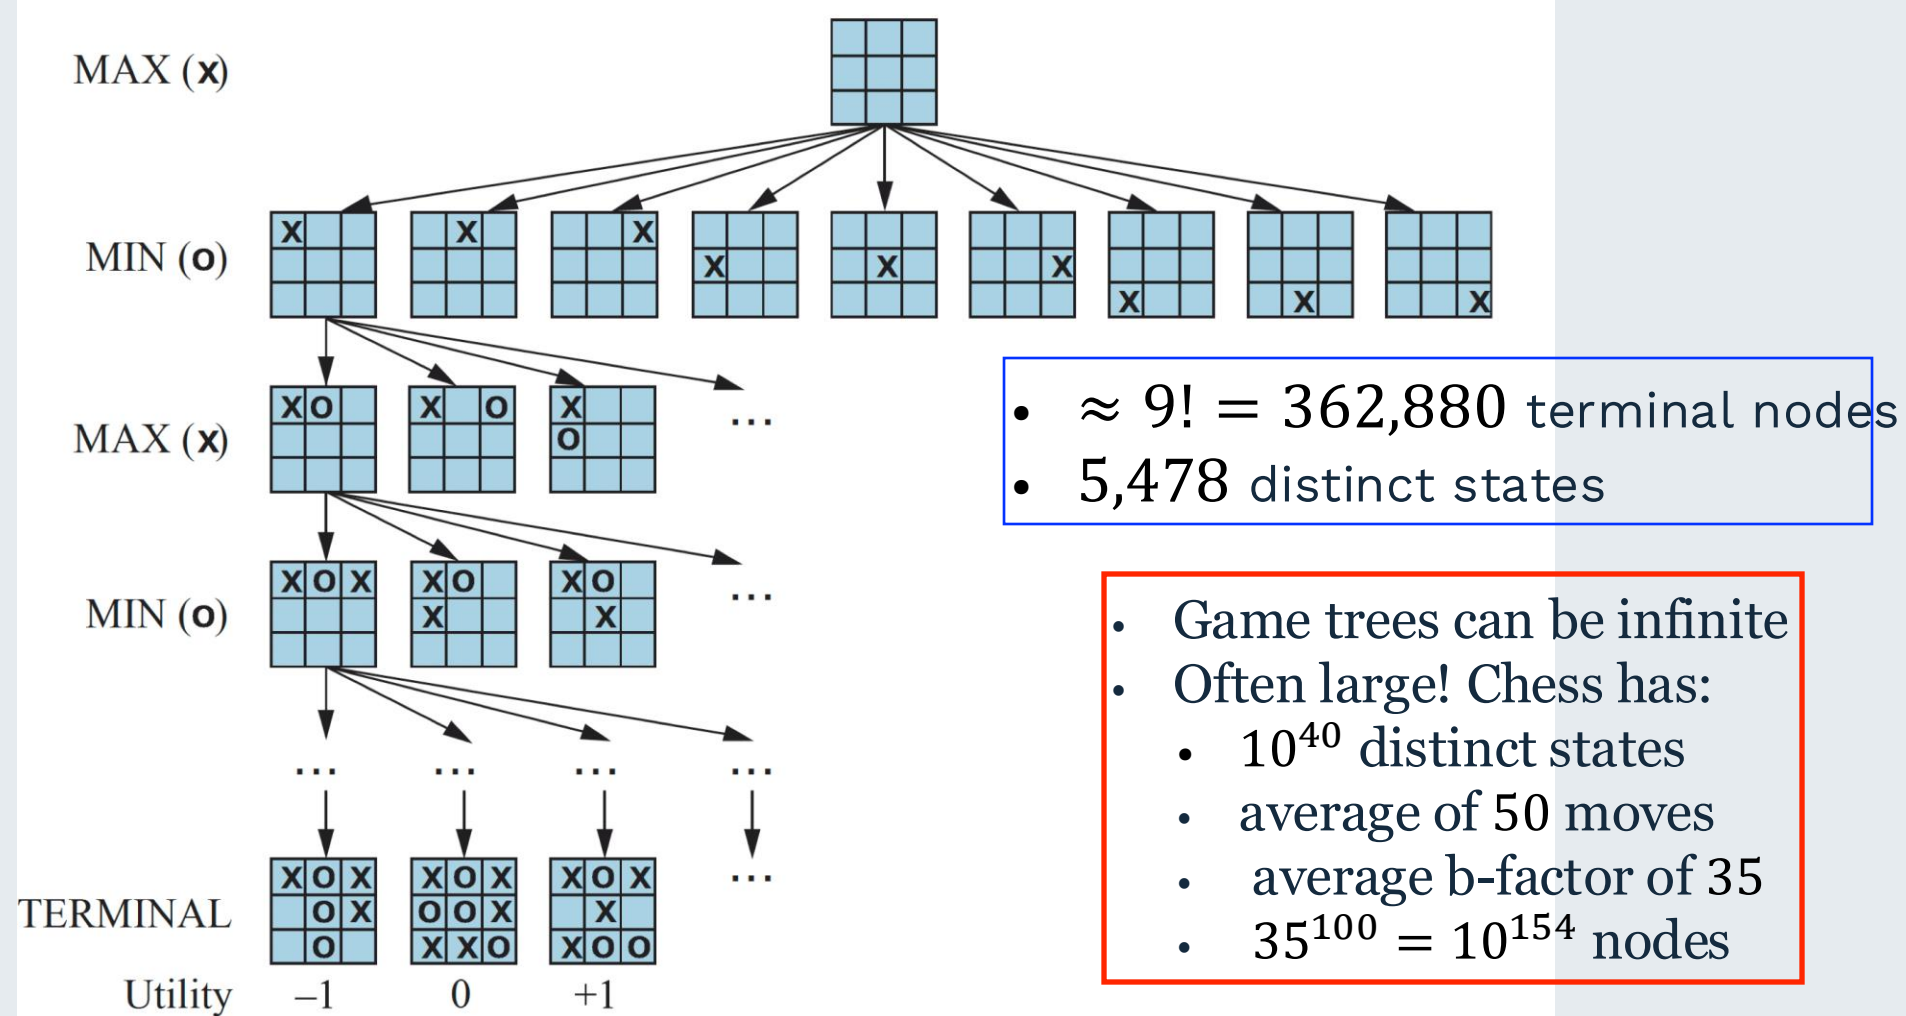
More generally: Adversarial Search

- Multi-Agent Environments
 - agents must consider the actions of other agents and how these agents affect or constrain their own actions.
 - environments can be **cooperative** or **competitive**.
 - One can view this interaction as a “game” and if the agents are competitive, their search strategies may be viewed as “adversarial”.
- Most often studied: **Two-agent, zero-sum games of perfect information**
 - Each player has a complete and perfect model of the environment and of its own and other agents actions and effects
 - Each player moves until one wins and the other loses, or there is a draw.
 - The utility values at the end of the game are always equal and opposite, thus the name zero-sum.
 - Chess, checkers, Go, Backgammon (uncertainty)

Games as Search

- The Game
 - Two players: One called MIN, the other MAX. MAX moves first.
 - Each player takes an alternate turn until the game is over.
 - At the end of the game points are awarded to the winner, penalties to the loser.
- Formal Problem Definition:
 - Initial State: S_0 – Initial board position
 - **TO-MOVE(s)** – The player whose turn it is to move in state **s**
 - **ACTION(s)** – The set of legal moves in state **s**
 - **RESULT(s,a)** – The transition model: the state resulting from taking action **a** in state **s**.
 - **IS-TERMINAL(s)** – A terminal test. True when game is over.
 - **UTILITY(s,p)** – A utility function. Gives final numeric value to player **p** when the game ends in terminal state **s**.
 - For example, in Chess: win (1), lose (-1), draw (0):

(Partial) Game Tree for Tic-Tac-Toe




Optimal Decisions in Games: Minimax Search

1. Generate the complete game tree using depth-first search.
2. Apply the utility function to each terminal state.
3. Beginning with the terminal states, determine the utility of the predecessor nodes (parent nodes) as follows:
 1. Node is a MIN-node
Value is the **minimum** of the successor nodes
 2. Node is a MAX-node
Value is the **maximum** of the successor nodes
4. From the initial state (root of the game tree), MAX chooses the move that leads to the highest value (**minimax** decision).

Note: Minimax assumes that MIN plays perfectly. Every weakness (i.e. every mistake MIN makes) can only improve the result for MAX.


Minimax Tree



- Interpreted from MAX's perspective
- Assumption is that MIN plays optimally
- The minimax value of a node is the utility for MAX
- MAX prefers to move to a state of maximum value and MIN prefers minimum value

What move should MAX make from the Initial state?

MAX utility values



$MINIMAX(s) =$

$$\begin{cases} UTILITY(s, MAX) & \text{if } IS-TERMINAL(s) \\ \max_{a \in Actions(s)} MINIMAX(RESULT(s, a)) & \text{if } TO-MOVE(s) = MAX \\ \min_{a \in Actions(s)} MINIMAX(RESULT(s, a)) & \text{if } TO-MOVE(s) = MIN \end{cases}$$

Minimax Algorithm

```

function MINIMAX-SEARCH(game, state) returns an action
  player  $\leftarrow$  game.TO-MOVE(state)
  value, move  $\leftarrow$  MAX-VALUE(game, state)
  return move

function MAX-VALUE(game, state) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
  v  $\leftarrow -\infty$ 
  for each a in game.ACTIONS(state) do
    v2, a2  $\leftarrow$  MIN-VALUE(game, game.RESULT(state, a))
    if v2 > v then
      v, move  $\leftarrow$  v2, a
  return v, move

function MIN-VALUE(game, state) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
  v  $\leftarrow +\infty$ 
  for each a in game.ACTIONS(state) do
    v2, a2  $\leftarrow$  MAX-VALUE(game, game.RESULT(state, a))
    if v2 < v then
      v, move  $\leftarrow$  v2, a
  return v, move

```

Assume max depth of the tree is m
and b legal moves at each point:

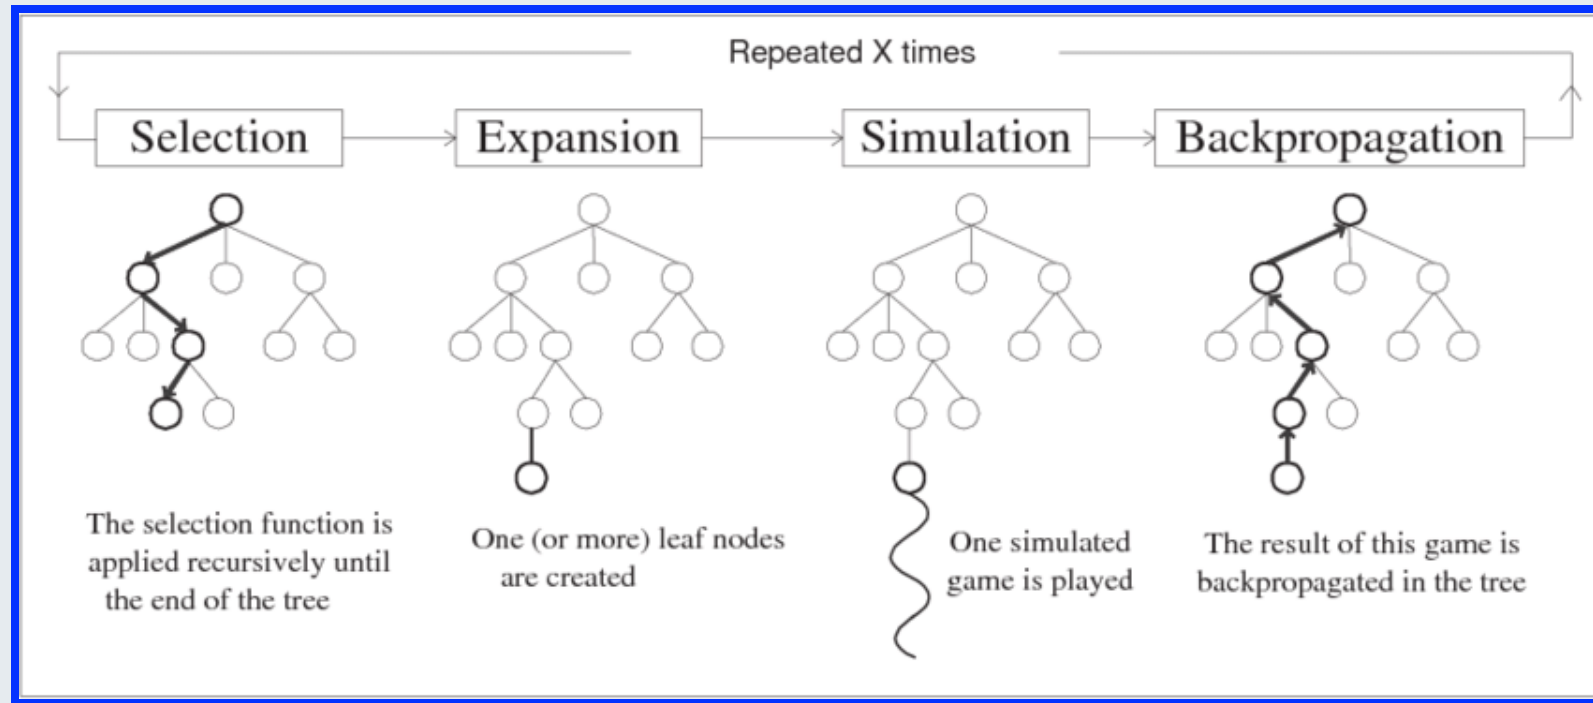
- Time complexity: $O(b^m)$
- Space complexity:
 - Actions generated at same time: $O(bm)$
 - Actions generated one at a time: $O(m)$

Serves as a basis for mathematical analysis
of games and development of approximations
to the minimax algorithm

Recursive algorithm that proceeds all the way down to the
leaves of the tree and then backs up the minimax values
through the tree as the recursion unwinds

4 Steps in MCTS

MCTS maintains a search tree and grows it on each iteration using the following steps:



Starting at the root of the search tree, choose a move using the [selection policy](#), repeating the process until a leaf node is reached

Grow the search tree by generating a new child/children.

Perform a [playout](#) from a child using the [playout policy](#). These moves are not recorded in the search tree

Use the simulation result to update the utilities of the nodes going back up to the root.

After X times: Choose the best move from start state